

Generative AI Leader Sample Questions

The Generative AI Leader sample questions are designed to help you familiarize yourself with the format of exam questions and example content that may be covered on the exam. These questions are for the current version of the exam that launched May 14, 2025.

Performance on the sample questions should not be used to predict your Generative AI Leader exam result.

Registration

First Name *

Ramoji Rao

Last Name *

Yalamati

Primary Email *

ramoji.rao.yalamati@accenture.com



Organization (Employer or School) *

Accentrue

Primary Relationship to Google *

Dropdown

Partner

Send me offers, updates and useful tips for getting the most out of Google Cloud training and certification products and services. *

Dropdown

Yes



- ✓ A human resources department deploys a generative AI (gen AI) model to screen job applications and provide a shortlist of candidates to recruiters. Recruiters notice that some seemingly qualified candidates are consistently being overlooked, but the AI provides no explanation for its rankings or exclusions. The company needs to address this lack of transparency. What should they do?
- ☐ Collect a larger and more diverse dataset for the gen AI model.
 - ☐ Fine-tune the gen AI model.
 - ☒ Implement explainable gen AI policies. ✓
 - ☐ Develop fairness assessments for the gen AI model.

Feedback

A is incorrect because, while a larger and more diverse dataset could potentially improve the model's accuracy and reduce bias, it does not inherently make the model's screening process transparent. Recruiters would still lack insight into why certain candidates are being shortlisted over others. While data quality is essential for ethical AI, it doesn't automatically provide explainability.

B is incorrect because fine-tuning aims to optimize the model's performance on the screening task, but it doesn't necessarily reveal how the model is making its choices. Even if the shortlist improves after fine-tuning, the lack of transparency regarding the selection criteria remains. Fine-tuning focuses on improving task-specific performance, not on making the model's reasoning understandable.

C is correct because it directly tackles the issue of unclear decision-making. Implementing explainable gen AI techniques would allow the HR department and recruiters to understand the factors that the AI model is using to rank and exclude candidates. This aligns with the principle that transparency is paramount for ethical applications and that users need to understand how AI systems make decisions. Explainable AI makes the decision-making processes of AI models transparent and understandable, which is crucial for building trust and identifying potential issues.

D is incorrect because, while developing fairness assessments is important for ensuring the AI model is not discriminatory in its candidate selection, it does not directly address the need to understand why the model is making specific decisions. Assessing for fairness helps identify potential biases, but it doesn't explain the reasoning behind individual candidate rankings or exclusions.



- ✓ A company is evaluating the use of large language models (LLMs) to enhance its operations and customer interactions. What is a primary characteristic of LLMs? *
- ☐ LLMs excel in highly specific technical tasks requiring deep, singular domain expertise.
 - ☐ LLMs learn and generalize effectively from small datasets for niche applications.
 - ☐ LLMs have strong inherent logical reasoning and problem-solving abilities without extra prompting.
 - ☒ LLMs are trained on vast datasets, enabling broad language and context understanding, and adaptability across many tasks. ✓

Feedback

A is incorrect because, while LLMs can be applied to specific tasks, their core strength lies in their general language understanding derived from broad training, enabling them to handle a variety of applications rather than being limited to highly specialized areas.

B is incorrect because LLMs are known for requiring vast amounts of data for effective training to achieve their broad capabilities. They are not typically effective with limited datasets without fine-tuning on more specific data.

C is incorrect because, while LLMs can exhibit reasoning capabilities, this often requires effective prompt engineering or advanced techniques like chain-of-thought prompting. They do not inherently possess strong logical reasoning without guidance. Additionally, limitations such as "hallucinations," indicating they are not flawless problem-solvers.

D is correct because foundation models, including LLMs, are trained on massive amounts of diverse data (text, images, and code). This broad training allows them to develop a deep understanding of the data and be adapted to many different tasks. The flexibility of foundation models, including LLMs, to support a wide range of use cases stems directly from this training on diverse data.



- ✓ An AI robot learns optimal package delivery routes in a city. It receives positive scores for fast, successful deliveries and negative scores for delays or failures. Through this feedback, the robot improves its navigation over time. What type of machine learning is being used to train the robot? *
- ☐ Supervised learning
 - ☐ Deep learning
 - ☐ Unsupervised learning
 - ☒ Reinforcement learning ✓

Feedback

A is incorrect because supervised learning requires labeled data, where each input is paired with a correct output. This scenario does not involve pre-labeled optimal routes; the robot learns through trial and error.

B is incorrect because, while deep learning can be a component of reinforcement learning, it is not the overarching learning paradigm described. Deep learning refers to neural networks with multiple layers. Deep learning can be used to implement the function that learns the optimal policy in reinforcement learning. However, the learning process itself is defined by the reward and penalty mechanism.

C is incorrect because unsupervised learning focuses on finding patterns in unlabeled data without specific goals or feedback. The robot's learning is goal-oriented (optimizing delivery routes) and driven by the feedback it receives.

D is correct because the robot learns through interaction with its environment and by receiving rewards (positive scores) and penalties (negative scores) for its actions. That is the fundamental process of reinforcement learning. Reinforcement learning is all about learning through interaction and feedback, where an algorithm learns to maximize rewards and minimize penalties by interacting with its environment.



- ✓ A company wants to use generative AI (gen AI) to automate complex workflows and improve decision-making across its various departments. They are considering implementing AI agents as a key component of their strategy. What is the primary function of an AI agent in a gen AI system? *
- ☐ To provide the computing power for training and running advanced AI models.
 - ☐ To be the user interface for interacting with AI models.
 - ☒ To be a smart system that can analyze, use tools, and make decisions to reach goals. ✓
 - ☐ To be a central storage place for the data that AI models use.

Feedback

A is incorrect because, while infrastructure is crucial for AI, its primary function is to provide the foundational computing resources, not to act as the intelligent decision-making component of an agent.

B is incorrect because, while platforms like Vertex AI Studio offer interfaces for interacting with models, the primary function of an agent is the autonomous reasoning and action, not just the user interface. Agents use models, but they are more than just an interface.

C is correct because AI agents are defined as entities that observe, act, and achieve goals by incorporating a reasoning loop and tools. A gen AI agent is an application that tries to achieve a goal by observing the world and acting upon it using the tools it has at its disposal. This highlights their ability to analyze situations and make decisions to fulfill objectives without constant human oversight. They leverage tools to interact with their environment and perform actions.

D is incorrect because, while data storage is essential for AI, it is a separate function from the agent's ability to process information, reason, and take actions. Agents utilize data stores as tools, but their core function is not data management itself.



✓ An advertising agency needs to quickly generate many photorealistic images from text for client campaigns because traditional photoshoots are slow and costly. They want to rapidly create high-quality visuals from text and reduce expenses. Which Google foundation model should they use? *

☐ Gemini

☐ Gemma

☐ Veo

☒ Imagen



Feedback

A is incorrect because, while Gemini is a multimodal model capable of understanding images and text, its primary strength for this specific scenario isn't solely focused on generating high-quality images from text as Imagen is.

B is incorrect because Gemma is a family of lightweight, open models built upon the research behind Gemini, but it is not specifically highlighted as the primary Google model for high-quality text-to-image generation.

C is incorrect because Veo is a model designed for generating video content from text or still images, not primarily for creating still photorealistic images from text descriptions.

D is correct because Imagen is a powerful text-to-image diffusion model that excels at generating high-quality images from textual descriptions. This directly addresses the agency's need to create photorealistic images from text prompts.



✓ A company is planning to integrate generative AI into its operations but is wary of becoming dependent on a single technology provider. They prioritize the ability to choose and integrate different AI tools and platforms as their needs evolve. Which inherent characteristic of Google Cloud would address this concern? *

- ☒ Google Cloud's emphasis on an open approach within its AI offerings. ✓
- ☐ Google Cloud's commitment to tightly integrated, proprietary AI solutions
- ☐ Google Cloud's strategy prioritizing fully managed AI services that simplify the user experience
- ☐ Google Cloud's primary focus on automating AI workflows

Feedback

A is correct because Google Cloud has an open approach and recognize the benefits of Google Cloud's open approach. This openness implies support for customer choice across different offerings. This directly mitigates concerns about vendor lock-in and promotes flexibility in their technology stack. Open standards allow users to move services between vendors more easily.

B is incorrect because a strict adherence to proprietary technologies would exacerbate concerns about vendor lock-in, directly contradicting the company's priority for flexibility.

C is incorrect because, while fully managed services offer convenience, they can potentially limit user control and choice. This would not align with the company's desire for flexibility and avoiding vendor lock-in.

D is incorrect because, while automation can streamline workflows, it doesn't inherently address the concern of vendor lock-in. The company's priority is maintaining flexibility and choice among different AI technologies.



- ✓ A consulting research team needs to analyze multiple lengthy reports and documents to find key trends and make client recommendations. They require a method to quickly understand each document's core findings, link information across sources, and efficiently organize insights for their report. Manual methods are too slow and complex. Which Google Cloud offering should they use? *

- ☒ NotebookLM ✓
- ☐ Gemini app
- ☐ Vertex AI Search
- ☐ Gemini for Google Workspace

Feedback

A is correct because NotebookLM is specifically designed as an AI-first notebook grounded in user-provided documents to help users gain insights faster. It allows users to upload multiple documents, ask questions about the content, request summaries, and save key insights as notes. This directly addresses the research team's need to analyze several documents, understand their findings, identify connections, and organize information efficiently. NotebookLM's focus on source-based answers ensures accuracy and the ability to trace back to the original information.

B is incorrect. The Gemini app is Google's generative AI chatbot that can help with writing, planning, learning, and more. However, it is not specifically designed for the in-depth analysis and organization of insights from a specific set of uploaded research documents in the way NotebookLM is.

C is incorrect because Vertex AI Search is designed for building search applications over structured and unstructured data. While it can help find information within documents, it does not offer the focused Q&A, summarization, and note-taking capabilities of NotebookLM for a specific research task.

D is incorrect because Gemini for Google Workspace integrates generative AI into productivity tools like Docs and Drive. While helpful for individual document-level tasks, it does not provide the dedicated, multi-document analysis and insight organization features of NotebookLM for a research project involving multiple complex documents.



✗ A grocery store chain has data in several internal systems like sales, inventory, and marketing. Employees waste time searching these systems for information on product performance, stock, and campaign effectiveness. They need a central way to easily access and understand data across these systems for better decisions and efficiency. Which Google Cloud offering should they use? *

☒ Gemini for Google Workspace

✗

☐ Google Agentspace

☐ Vertex AI Search

☐ Conversational Agents

Correct answer

☒ Google Agentspace

Feedback

A is incorrect because Gemini for Google Workspace integrates generative AI features into productivity tools like Gmail, Docs, and Sheets to enhance individual and collaborative work within those applications. While helpful for tasks within these tools, it doesn't provide a centralized solution for accessing information across multiple disparate internal systems like Agentspace.

B is correct because Google Agentspace is designed to help teams use their company's information more effectively by creating customized agents that can access and understand data from various sources, regardless of where that data is stored. This directly addresses the grocery store's need for a central platform where employees can easily find information across their sales, inventory, and marketing systems to improve decision-making and operational efficiency.

C is incorrect because Vertex AI Search is designed for building search applications over structured and unstructured data to help users find specific information. While it could potentially search across the grocery store's systems, it doesn't offer the broader capabilities of Google Agentspace to create customized AI assistants that can proactively provide insights and automate information retrieval from various systems.

D is incorrect because Conversational Agents are primarily used for building customer-facing chatbots to automate interactions and answer customer inquiries. While valuable for customer service, they are not the most suitable solution for improving internal employee access to information across various internal systems, which is the core need in this scenario.



- ✓ A tech company has separate teams using different tools for their machine learning projects, causing duplicated work and scaling issues. They need a central platform to manage all their AI development, deployment, and monitoring efficiently. Which Google Cloud offering should they use? *

- ☐ Cloud Functions
- ☒ Vertex AI
- ☐ Google Agentspace
- ☐ BigQuery



Feedback

A is incorrect because Cloud Functions is a serverless execution environment for running event-driven code. While Cloud Functions can be a component in an AI workflow, it does not provide a central platform to manage the entire machine learning lifecycle across teams.

B is correct because Vertex AI is Google Cloud's unified ML platform designed to streamline the entire ML workflow. Vertex AI offers the infrastructure, tools, and pre-trained models needed for the centralized building, deployment, and management of ML and generative AI solutions. Vertex AI encompasses data preparation, model training, evaluation, deployment, and monitoring, facilitating collaboration and efficient resource use for diverse AI initiatives.

C is incorrect because Google Agentspace is designed to help teams use their company's information more effectively by creating customized AI agents for information access and task automation. While it can leverage ML models, it is not the primary platform for end-to-end ML project management.

D is incorrect because BigQuery is a fully managed and serverless data warehouse optimized for scalable data analysis. It is crucial for storing and analyzing large datasets used in machine learning and integrable with Vertex AI. However, it doesn't offer the comprehensive platform for managing the complete ML development, deployment, and monitoring process that Vertex AI does.



✓ A software company's AI chatbot struggles to answer customer questions about recently released features because this information is not in its original training data. Customers are getting inaccurate answers, increasing support agent workload. The company wants the chatbot to use the latest product documentation to give accurate, up-to-date responses without retraining the entire model. Which technique should they use? *

- ☐ Fine-tuning
- ☐ Prompt engineering
- ☒ Retrieval-augmented generation (RAG) ✓
- ☐ Human-in-the-loop (HITL)

Feedback

A is incorrect because fine-tuning requires retraining the model on new data, which is more time-consuming than immediately accessing existing documentation with RAG.

B is incorrect because prompt engineering alone cannot provide information the model doesn't already have in its training data. The new feature details are likely outside this knowledge base. However, prompt engineering can be used within a RAG system.

C is correct because retrieval-augmented generation (RAG) allows the language model to retrieve relevant information from external sources, like the latest product documentation, and use it to generate more accurate and contextually appropriate responses, directly addressing the chatbot's knowledge gap about new features without full retraining.

D is incorrect because human-in-the-Loop (HITL) involves human intervention and doesn't proactively enable the chatbot to answer questions about new information automatically as RAG does.



- ✓ A business analyst asks a generative AI model about the quarterly revenue of a small startup that recently entered the market. The model confidently provides a specific revenue figure and even mentions a supposed press release detailing the company's success. However, after further investigation, the analyst discovers that the startup has not yet released any financial reports, and no such press release exists. The information provided by the AI model is entirely fabricated despite sounding plausible. Which type of large language model limitation does this exemplify? *

- ☐ Bias
- ☐ Knowledge cutoff
- ☐ Data dependency
- ☒ Hallucinations



Feedback

A is incorrect because bias refers to the model's tendency to produce outputs that reflect imbalances or prejudices present in its training data, not necessarily the fabrication of entirely new information.

B is incorrect because knowledge cutoff refers to the point in time after which the model has not been trained on new information. While this can lead to an inability to answer questions about recent events, in this scenario, the model is generating non-existent information rather than stating it doesn't know or providing outdated information.

C is incorrect because data dependency highlights that the performance of foundation models relies heavily on the quality and completeness of their training data. While a lack of specific data might lead to an inability to answer, hallucinations involve generating factually incorrect or nonsensical responses, even if the model believes it has the information.

D is correct because hallucinations occur when foundation models produce outputs that are not accurate or based on real information. In this scenario, the AI model fabricates revenue data and a press release that does not exist, which is a clear example of a hallucination. The model cannot verify information against external sources and may generate convincing but incorrect responses.



- ✓ A generative AI tool that answers employee policy questions is providing outdated and inaccurate information, causing confusion. The company wants the tool to give reliable answers based on the latest official documents. What should the organization do? *
- ☐ Fine-tune the underlying language model with a broader dataset of general knowledge.
 - ☐ Increase the temperature setting of the language model.
 - ☒ Implement grounding techniques. ✓
 - ☐ Reduce the token count parameter.

Feedback

A is incorrect because fine-tuning with a broader general knowledge dataset would not necessarily ensure the AI adheres to the company's specific and latest internal policies.

B is incorrect because increasing the temperature parameter would likely increase inaccurate responses .

C is correct because grounding connects the AI's output to verifiable information sources, such as internal documents, improving accuracy and reliability. Retrieval-augmented generation (RAG) is a specific grounding method that retrieves relevant information before generating a response.

D is incorrect because reducing the token count only affects the length of the responses, not their accuracy.



✓ A sales team wants to create dynamic and personalized video pitches for potential clients. They receive client information in various formats and need an AI model that can transform this information into engaging video content tailored to each client's specific needs and challenges. Which Google model should they use? *

☐ Gemma

☐ Gemini

☐ Imagen

☒ Veo



Feedback

A is incorrect because Gemma is a family of lightweight, open models primarily focused on developer-friendly and customizable solutions. While versatile, its core strength isn't video generation for sales pitches.

B is incorrect because Gemini is a multimodal model capable of understanding and operating across diverse data formats like text, images, and audio for various tasks. However, its primary function is not specifically video content creation for sales.

C is incorrect because Imagen specializes in generating high-quality images from textual descriptions, which does not directly address the team's need for video pitch creation.

D is correct because Veo is a model capable of generating video content from text or still images. This directly addresses the sales team's need to transform client information into engaging and personalized video pitches.



✓ A growing retail company with fragmented phone, email, and basic website chat support needs a unified cloud solution. They require integrated communication channels, consistent customer experiences, and scalable support that ensures security and privacy. Which Google Cloud offering should they use? *

- ☐ Vertex AI Platform
- ☒ Google Cloud Contact Center as a Service ✓
- ☐ Conversational AI
- ☐ Vertex AI Search

Feedback

A is incorrect because Vertex AI Platform is Google Cloud's unified machine learning platform for building, deploying, and managing ML models. While it can be used to enhance customer experience through AI-powered features, it does not provide a complete contact center solution with integrated communication channels and agent management.

B is correct because Google Cloud Contact Center as a Service (CCaaS) offers a complete cloud-based contact center solution. It integrates channels (phone, text, email, etc.), ensures consistent experiences (omnichannel support), and provides scalable and secure support. It manages the infrastructure, integrates with CRMs, and handles agent routing.

C is incorrect because Conversational AI is a broader term referring to the technology used to build interactive systems like chatbots. While Google Cloud Contact Center as a Service leverages Conversational AI in its tools like Conversational Agents, option B does not represent a full contact center infrastructure offering the required integration and scalability.

D is incorrect because Vertex AI Search is a Google Cloud service that allows developers to integrate advanced search capabilities into their applications and handle large volumes of multimodal data. While it can improve self-service and help customers find information, it does not provide the integrated communication channels, agent management, and comprehensive features required for a unified contact center solution.



- ✓ An organization is seeking to improve how its employees access and use internal company information scattered across various systems. They want to provide their knowledge workers with tools that can understand and use this data to enhance productivity and decision-making. What is a key benefit of using Google Cloud Agentspace in this scenario? *
- ☐ Agentspace primarily focuses on enhancing external customer engagement through AI-powered chatbots.
 - ☐ Agentspace directly manages the underlying infrastructure and hardware required for AI model training.
 - ☒ Agentspace allows employees to find and use internal information more easily by creating custom AI agents that can access and understand data from various enterprise sources. ✓
 - ☐ Agentspace is mainly designed for building and deploying custom machine learning models for predictive analytics.

Feedback

A is incorrect because, while Google offers tools for external customer engagement, Google Agentspace is specifically designed to help internal teams use their company's information more effectively.

B is incorrect because, while Google Cloud provides AI-optimized infrastructure, Google Agentspace is a platform built on top of this infrastructure. It focuses on the creation and deployment of AI agents, not the direct management of hardware.

C is correct because Google Agentspace uses AI to create customized agents that can access and understand data from various sources within an organization, regardless of where it is stored. This allows employees to more easily find relevant information, conduct research, and automate tasks, ultimately increasing productivity.

D is incorrect because, while Vertex AI Platform is Google Cloud's unified machine learning platform for building, deploying, and managing models, Google Agentspace builds upon these capabilities by providing a platform specifically for deploying AI agents that leverage enterprise data for knowledge work.

Untitled Section



✓ What is reinforcement learning? *

- ☐ Learning from labeled data with correct output pairs.
- ☐ Learning by identifying patterns in unlabeled data.
- ☒ Learning through interaction and feedback. ✓
- ☐ Learning by training on vast data to generate new content.

Feedback

A is incorrect because this definition describes supervised learning.

B is incorrect because this definition describes unsupervised learning.

C is correct because reinforcement learning is accurately defined as a process where an agent learns by interacting with an environment and receiving feedback in the form of rewards or penalties.

D is incorrect because this describes the process to train generative AI models. This definition does not explicitly define reinforcement learning, but understanding its core mechanism of learning through interaction is a fundamental concept in machine learning.

Untitled Section



- ✓ A company is developing a system to automatically categorize customer support emails. They have a collection of thousands of past emails, and each email has been manually reviewed and tagged with a category such as "Billing Inquiry," "Technical Support," or "Feature Request." What type of data is this? *

- ☐ Unlabeled data
- ☒ Labeled data ✓
- ☐ Structured data
- ☐ Raw data

Feedback

A is incorrect because unlabeled data is simply data that is not tagged or labeled in any way. It's raw, unprocessed information without inherent meaning. The emails in the stem have been processed and assigned categories, thus they are not unlabeled.

B is correct because these categories; Billing Inquiry, Technical Support, or Feature Request; function as labels, which assigns meaning to each email by providing a corresponding output. This pairing of input (email) and output (label) defines labeled data, which is crucial for supervised machine learning models that learn from such input-output pairs.

C is incorrect because structured data has information neatly arranged in tables and is easy to search and find. Examples include online shopping orders or bank statements. While the email categories might be considered structured information, the content of the emails themselves (the text) is likely unstructured data as it lacks a predefined format and isn't easily organized into rows and columns

D is incorrect because raw data refers to data that has not been processed or formatted for use. While the original emails might have started as raw data, the process of manually reviewing and tagging them with categories has transformed them into labeled data. The labels provide a level of processing and meaning that distinguishes them from unprocessed raw data.

Untitled Section



✓ What is the definition of a generative AI (gen AI) model? *

- ☐ A physical device that houses the hardware components of a gen AI system.
- ☒ A complex algorithm trained on vast amounts of data to learn patterns and relationships. ✓
- ☐ A user interface that allows users to interact with a gen AI system.
- ☐ A set of rules and guidelines governing responsible development and use of gen AI.

Feedback

A is incorrect because a gen AI model is not a physical device.

B is correct because a generative artificial intelligence (gen AI) model operates as a sophisticated algorithm. This algorithm undergoes an extensive training process, being exposed to and learning from remarkably large and diverse datasets. Through this process, the model identifies intricate patterns, subtle nuances, and complex interrelationships present within the data. This learned understanding of the underlying structure and statistical distributions of the training data then empowers the gen AI model to generate novel and original content that reflects the characteristics of the data it was trained on. This generated content can take various forms, including text, images, audio, or other types of data, depending on the model's architecture and training objectives.

C is incorrect because a gen AI model is not a user interface.

D is incorrect because a gen AI model is not a set of rules and guidelines.

Untitled Section



✗ A video game company created a virtual reality game with virtual characters that can interact with users in a more natural and intuitive way by using gestures and facial expressions to communicate. What type of agent is this? *

- ☐ Creative agent
- ☐ Workflow agent
- ☒ Virtual assistant agent ✗
- ☐ Conversational agent

Correct answer

- ☒ Conversational agent

Feedback

A is incorrect because while creative agents can be used in games, the primary function described here aligns more with a conversational agent.

B is incorrect because a workflow agent focuses on automating tasks, not natural interaction.

C is incorrect because a virtual assistant typically provides support and information, not the interactive dialogue described.

D is correct because a conversational agent is designed to understand and respond in a way that makes sense, aligning with the described interaction.

Untitled Section



✓ A software company has developers who need to write, review, debug, and generate code from natural language descriptions by using generative AI. What type of agent is this? *

- ☐ Data analysis agent
- ☐ Workflow agent
- ☐ Data agent
- ☒ Code agent



Feedback

A is incorrect because a data analysis agent supports data-related tasks and not specifically code tasks.

B is incorrect because a workflow agent focuses on automating tasks, not the interactive code assistance described.

C is incorrect because a data agent analyzes data, not code, to identify trends and insights.

D is correct because a code agent can assist developers in writing, reviewing, and debugging code, and even generating code from natural language descriptions.

Untitled Section



- ✓ An advertising agency needs to quickly create many different photorealistic * images from text descriptions for client campaigns because traditional photoshoots are too slow and costly. The agency's goal is to generate high-quality visuals rapidly from text to improve campaign speed and lower expenses. Which Google foundation model should be used?

- ☐ Gemini
- ☐ Gemma
- ☐ Veo
- ☒ Imagen



Feedback

A is incorrect because while Gemini is multimodal, Imagen is more specialized for photorealistic image generation from text.

B is incorrect because Gemma is focused on text and cannot create images.

C is incorrect because Veo generates videos, not the still images required by the agency.

D is correct because Imagen's main strength is generating high-quality images from text, which directly solves the agency's problem of needing fast and affordable visual content creation for advertising.

Untitled Section



✓ A research scientist wants to use Veo to visualize live, fluctuating data streams on a real-time dashboard. Why would Veo be a poor choice for this particular task? *

- ☒ Veo generates videos from static inputs like text or images and cannot process or dynamically visualize live data. ✓
- ☐ Veo may lack specific scientific visualization styles needed for accurate data representation.
- ☐ Veo use for this scenario would require too many computational resources.
- ☐ Veo is designed for short-form video, not continuous, long-duration live data displays.

Feedback

A is correct because Veo generates videos from static inputs (text, images) and cannot process or dynamically visualize live data. The research scientist's task requires a system capable of real-time processing and visualization of live, fluctuating data streams, which is fundamentally different from Veo's design.

B is incorrect because while specific scientific visualization styles are important, the primary reason Veo is unsuitable is its inability to handle live, fluctuating data streams. Its design for static inputs makes it incapable of real-time dynamic visualization, regardless of its graphical styles.

C is incorrect because the provided sources do not contain information about Veo's computational resource requirements. The core issue is the functional mismatch between Veo's design and the task's need for dynamic data processing.

D is incorrect because the sources do not specify limitations on video duration for Veo. Even if they did, the main problem remains Veo's fundamental inability to process and dynamically visualize live, fluctuating data in real-time, which is distinct from and more critical than potential output video length.

Untitled Section



✓ A company sells custom-designed phone cases on their website. How should Imagen be effectively used? *

- ☒ To generate realistic images of phone cases on devices from text descriptions of designs. ✓
- ☐ To analyze customer feedback to identify popular phone case design trends.
- ☐ To predict demand for different phone case designs based on sales data.
- ☐ To transcribe customer audio feedback on prototypes of phone case designs.

Feedback

A is correct because Imagen is trained to generate high-quality images from text descriptions. The company can use Imagen to create realistic visuals of their custom phone case designs by providing textual descriptions, which is ideal for product mockups on their website.

B describes text analysis, a task better suited for language models like Gemini, not Imagen's image generation capabilities.

C involves data analysis and prediction, which are typically handled by traditional machine learning models and analytical tools, not Imagen.

D describes audio processing and transcription, a task for which speech-to-text models like Chirp would be appropriate, not Imagen.

Untitled Section



✓ A company lacks extensive in-house machine learning and AI expertise. *

How does Google Cloud democratize AI?

- ☐ By offering exclusive access to its most powerful AI models for high-spending clients.
- ☐ By providing fully automated AI solutions requiring no technical user input.
- ☒ By providing a comprehensive AI platform with low-code/no-code tools, pre-trained models, and easy-to-use APIs. ✓
- ☐ By offering free custom AI solution development and deployment for all businesses.

Feedback

A is incorrect because it describes limiting access based on financial commitment, rather than making AI widely available. Democratization implies broader access for all.

B is incorrect because while Google Cloud offers tools that automate certain tasks and simplify workflows, they also emphasize the need for humans-in-the-loop in the AI implementation process. The platform offers different levels of technical input, from no-code/low-code to full code development, indicating that while some solutions require minimal technical knowledge, they are not universally "fully automated requiring no technical input."

Option C is correct because Google Cloud's AI platform democratizes AI development by offering low-code/no-code tools such as Vertex AI AutoML, access to a wide range of pre-trained models in Model Garden, and user-friendly APIs. This comprehensive approach empowers individuals with different technical backgrounds to leverage AI without needing deep machine learning expertise.

D is incorrect because this describes an unsustainable business model and circumvents the need for users to use the platform to build their own solutions, which is a core aspect of democratizing AI.

Untitled Section



- ✓ A marketing team is brainstorming creative ideas for a new social media campaign. They want to quickly generate various taglines and initial post drafts to explore different angles and messaging options. They need a Google Cloud prebuilt generative AI tool without additional setup that allows for rapid content creation and experimentation. What should the team do? *
- ☐ Use NotebookLM to upload past successful campaigns and analyze their key elements.
 - ☐ Use Gemini for Google Workspace within Google Slides to prepare a presentation outlining their social media campaign strategy.
 - ☒ Use the Gemini app to create taglines and social media post drafts that align with their campaign goals and target audience. ✓
 - ☐ Create a custom Gem in Gemini Advanced with specific brand guidelines and target audience profiles.

Feedback

A is incorrect because NotebookLM is primarily for analyzing uploaded documents and extracting insights, not for rapid, open-ended creative content generation for social media. B is incorrect because while Gemini for Google Workspace can help refine and present content, it is not the primary tool for initial, rapid generation of multiple creative taglines and post drafts.

C is correct because the Gemini app is designed for direct interaction and content generation through prompting. The marketing team can quickly input their campaign goals and target audience details and ask Gemini to generate multiple tagline options and initial post drafts for rapid brainstorming and exploration.

D is incorrect because while creating a custom Gem can ensure brand consistency, it requires more setup and is less suited for the initial, exploratory phase of generating a wide range of creative ideas quickly. The Gemini app offers a more immediate solution for this purpose.

This form was created inside Google.com. [Privacy & Terms](#)

Google Forms



